


# CSE7850/CX4803 Machine Learning in Computational Biology



## Lecture 9: Deep Learning for Sequence Data

Yunan Luo


Week	Date	Topic	Contents
1	01/08	Introduction	Introduction & Logistics
1	01/10	Basics in computational biology	Molecular biology
2	01/15		No class (MLK day)
2	01/17		Sequence alignment I
3	01/22		Sequence alignment II
3	01/24	ML foundations	No Class (PyTorch video + exercise)
4	01/29		Regression & Gradient descent
4	01/31		Classification & Toolbox for Applied ML
5	02/05		Neural networks
5	02/07		Deep learning
6	02/12	Learning from sequence data	Deep learning for Protein/DNA sequences
6	02/14		Large language models (LLMs)
7	02/19	Learning from high-dim data	Clustering and dimensionality reduction
7	02/21		Generative AI
8	02/26	Learning from network data	Network basics & ML for graphs
8	02/28		Graph neural network
9	03/04	Learning from structure data	Protein structure prediction & generation (AlphaFold, diffusion models)

# Today's plan

- Biological sequences
- Deep learning for sequence data in biomedicine
  - Supervised learning
    - CNN, RNN, LSTM, Transformer
  - Self-supervised learning
    - Overview of language modeling (LLM)

# Protein sequence

MFVFLVLLPLVSSQCVNLTTRTQLPPAYTNSFTRGVYYPDKVFRSSVLHSTQDLFLPFFS  
NVTWFHAIHVSGTNGTKRFDNVLFPFNDGVYFASTEKSNIIRGWIFGTTLDSKTQSLLIV  
NNATNVVIKVCEFQFCNDPFLGVYYHKNNKSWMSESEFRVYSSANNCTFEYVSQPFLMDLE  
GKQGNFKNLREFVFKNIDGYFKIYSKHTPINLVRDLPGGFSALEPLVDLPIGINITRFQT  
LLALHRSYLTPGDSSSGWTAGAAAYYVGYLQPRTFLLKYNENGTITDAVDCALDPLSETK  
CTLKSFTVEKGIYQTSNFRVQPTESIVRFPNITNLCPFGEVFNATRFASVYAWNRRKISN  
CVADYSVLVNSASFSTFKCYGVSPTKLNLDLCFTNVYADSFVIRGDEVQRQIAPGQTGKIAD  
YNYKLPPDDFTGCVIAWNSNNLDSKVGNNYNYLYRLFRKSNLKPFERDISTEIQAGSTPC  
NGVEGFNCYFPLQSYGFQPTNGVGYPYRVVLSFELLHAPATVCGPKKSTNLVKNKCVN  
FNFNGLTGTGVLTESNKKFLPFQQFGRDIADTTDAVRDPQTLIELDITPCSFGGVSVITP  
GTNTSNQVAVLYQDVNCTEVPVAIHADQLTPTWRVYSTGNSNVFQTRAGCLIGAEHVNNYSY  
ECDIPIGAGICASYQTQTSNPRRARSVASQSI IAYTMSLGAENSVAYSNNISIAIPTNFTI  
SVTTEILPVSMTKTSVDCTMYICGDS TECSNLLLQYGSFCTQLNRALTGIAVEQDKNTQE  
VFAQVKQIYKTPPIKDFGGFNFSQILPDPSKPSKRSFIEDLLFNKVTLADAGFIKQYGDC  
LGDIAARDLICAQKFNGLTVLPLLTDEMIAQYTSALLAGTITSGWTFGAGAAALQIPFAM  
QMAYRFNGIGVTQNVLYENQKLIANQFN SAIGKIQDSLSTASALGKLQDVVNQNAQALN  
TLVKQLSSNFGAISSVLNDILSRLDKVEAEVQIDRLITGRLQSLQTYVTQQLIRAAEIRA  
SANLAATKMSECVLGQSKRVDFCGKG YHLM SFPQSAPHGVVFLHVTYVPAQEKNTTAPA  
ICHDGKAHFPPREGVVFVSNGTHWFVTQRNFYEPQIITDNTFVSGNCDVVIGIVNNTVYDP  
LQPELDSFKEELDKYFKNHTSPDVLGDISGINASVVNIQKEIDRLNEVAKNLNESLIDL  
QELGKYEQYIKWPWYIWLGFIAGLIAIVMVTIMLCCMTSCCSC LKGCCSCGSCCKFDEDD



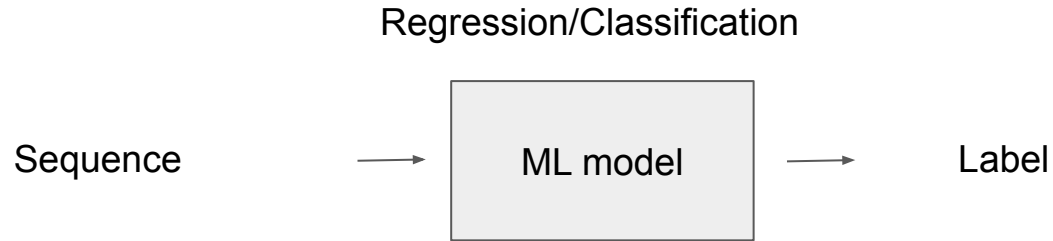
SARS-CoV-2 spike protein

# Summary: biological sequences

- DNA = nucleotide sequence
  - Alphabet size = 4 (A,C,G,T)
- RNA (single stranded)
  - Alphabet size = 4 (A,C,G,U)
- Protein sequence
  - Alphabet size = 20

# Deep learning for sequence data

# Problem formulation



Example: protein stability prediction

Input	Output
...DNGVDGEWTYDDATKTFTVTE	1.0
...DNGCDGEWTYDDATKTFTVTE	-0.2
...DNGVWGEWTYDDATKTFTVTE	3.9
...DNGVWGEWTYDDATKTFTFTE	5.4
...DNGVMGEWTYDDATKTFTDTE	-0.1

# Sequence encoding

- one-hot encoding


	C	G	A	T	A	A	C	C	G	A	T	A	T
A	0	0	1	0	1	1	0	0	0	1	0	1	0
C	1	0	0	0	0	0	1	1	0	0	0	0	0
G	0	1	0	0	0	0	0	0	1	0	0	0	0
T	0	0	0	1	0	0	0	0	0	0	1	0	1

- contextual embedding (language models)
  - Rives, Alexander, et al. "Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences." *Proceedings of the National Academy of Sciences* 118.15 (2021).




Model 1:  
Convolutional Neural Network (CNN)


# 1D convolution




# 1D convolution



# 1D convolution



# 1D convolution



Model 2:  
Recurrent Neural Network (RNN)

## Model #2: Recurrent Neural Network (RNN)

## Example: sequence labeling problems

- Part of speech



Article	Noun	Verb	Preposition	Article	Noun
<b>The</b>	<b>cat</b>	<b>sat</b>	<b>on</b>	<b>the</b>	<b>mat.</b>

- Handwriting recognition


Foreign Minister. → FOREIGN MINISTER.

- Protein secondary structure prediction

NKEILDEAYVMASVDNPHVCRLLGICLTSTVQLITQLMPFGCLLDYVREHKDNI GSQYLL



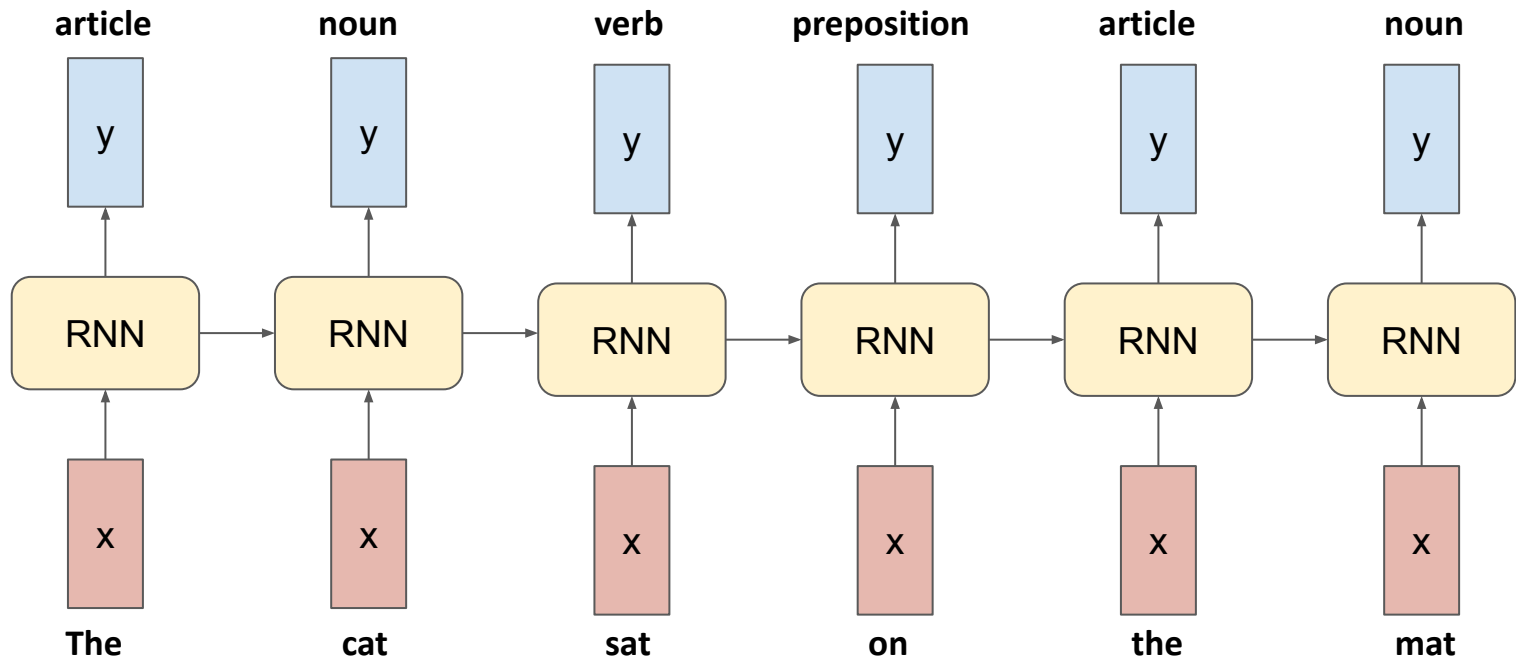
# Recurrent Neural Network (RNN)



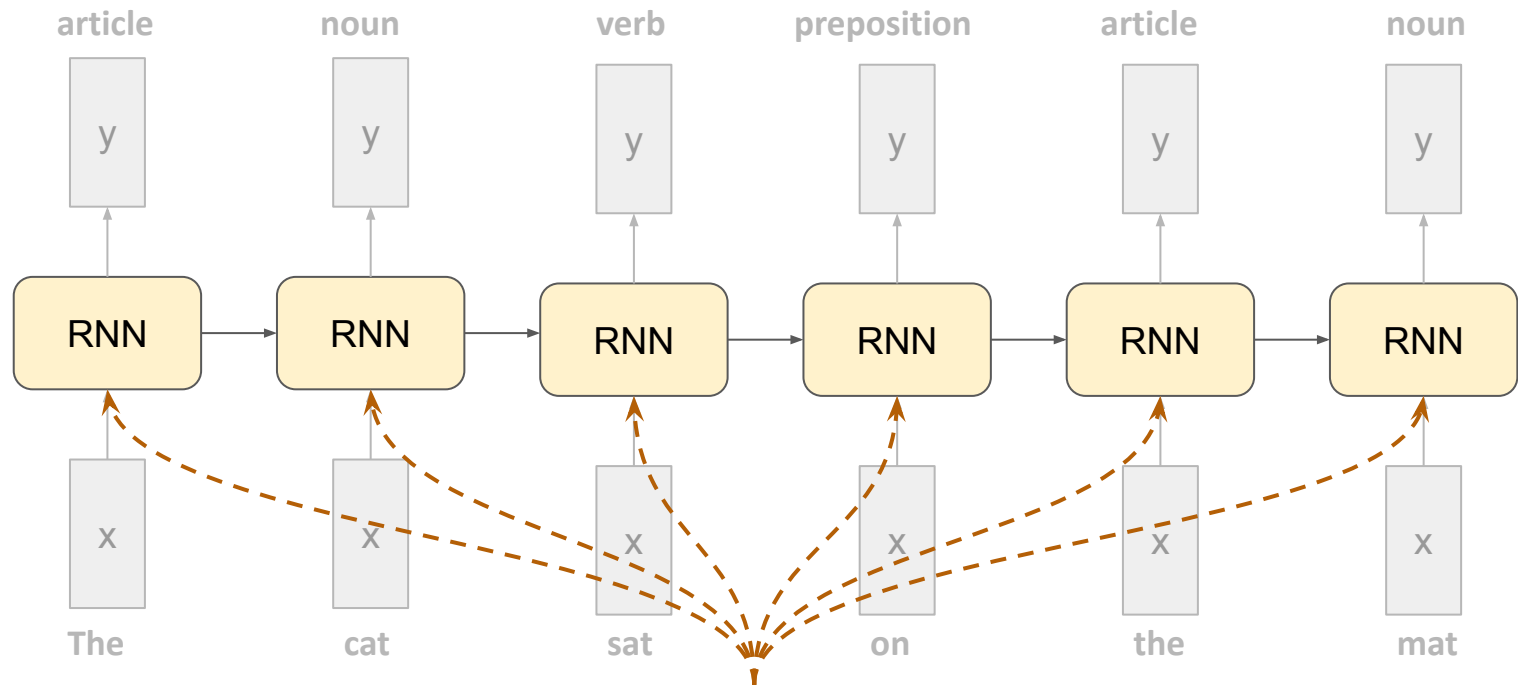
**Key idea:** RNNs have an “internal state” that is updated as the input sequence is processed



# Unrolled RNN



# Unrolled RNN




The same set of function and the same set of parameters are used at every time step

# RNNs for sequence processing

one to many


(e.g., image  
captioning)



# RNNs for sequence processing


one to many

(e.g., image captioning)



many to one


(e.g., protein function prediction)



# RNNs for sequence processing


one to many

(e.g., image captioning)




many to one

(e.g., protein function prediction)



many to many


(e.g., protein structure prediction)



# RNNs for sequence processing


one to many

(e.g., image captioning)




many to one

(e.g., protein function prediction)




many to many

(e.g., protein structure prediction)




many to many

(e.g., auto completion)



# RNN hidden state update


At every step  $t$ , the hidden state is updated based on the previous state and the current input



$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state                      function with  
parameter  $W$                       previous  
state                      Input vector  
at time step  $t$

# RNN output



At every step  $t$ , the output is generated based on the current state

Another function with  
parameter  $\theta$

$$\boxed{y_t} = \boxed{f_\theta}(\boxed{h_t})$$

new state                      new state




# Further readings of RNN

- *Deep Learning*, [Chapter 6](#)
- What is function  $f()$ ?
  - $f()$  is usually called “unit” in RNN
  - It defines a “computational graph” the produces  $h_t$  based on  $h_{t-1}$  and  $x_t$
- Popular RNN variants
  - Long short-term memory (LSTM)
  - Gated recurrent unit (GRU)

function with  
parameter  $W$


$$h_t = f_W(h_{t-1}, x_t)$$



([image source](#))

# Summary: key ideas of RNNs

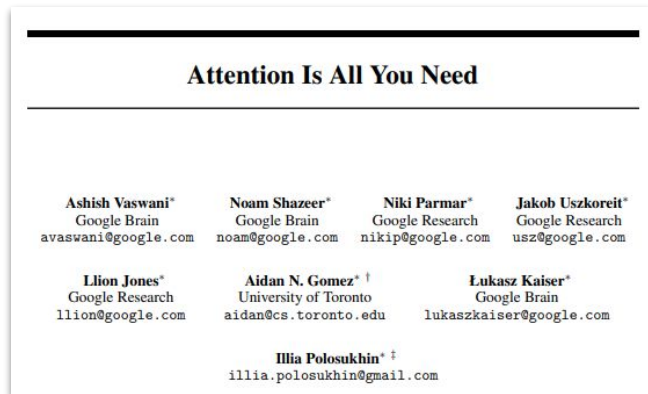
- Process sequence data with **variable lengths**
  - DNA/RNA/protein sequences, text, audio, time series data
- Capture **sequential** (temporal) information/dependencies in the data
- Parameters **shared** over time steps
- Common to use LSTM or GRU



is it ?

# Model 3: Transformers

# Model #3: Transformers



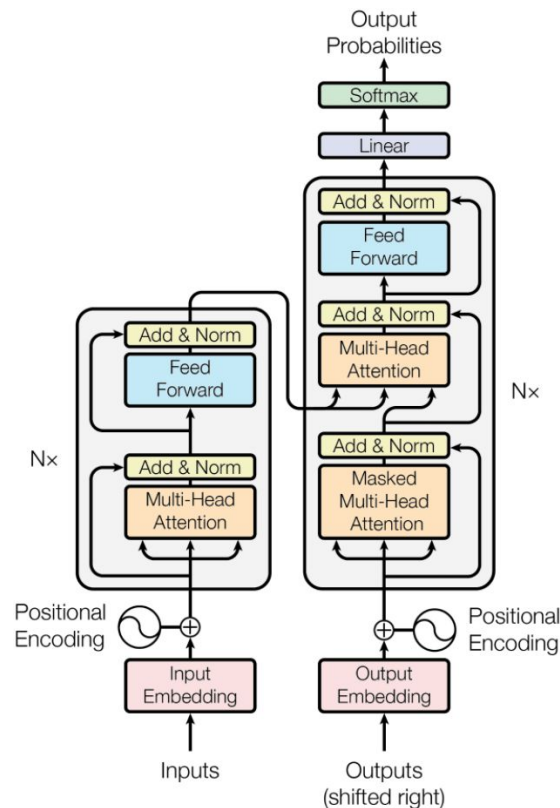
(NeurIPS 2017)

- Encoder-Decoder
- Sequence-to-sequence
- Transforms one sequence into another sequence, using full context of each

## Attention is all you need

A Vaswani, N Shazeer, N Parmar... - Advances in neur  
... to attend to **all** positions in the decoder up to and inc  
... We implement this inside of scaled dot-product **atten**  
☆ Save 77 Cite Cited by 108423 Related articles

(Citation as of 02/11/2024)



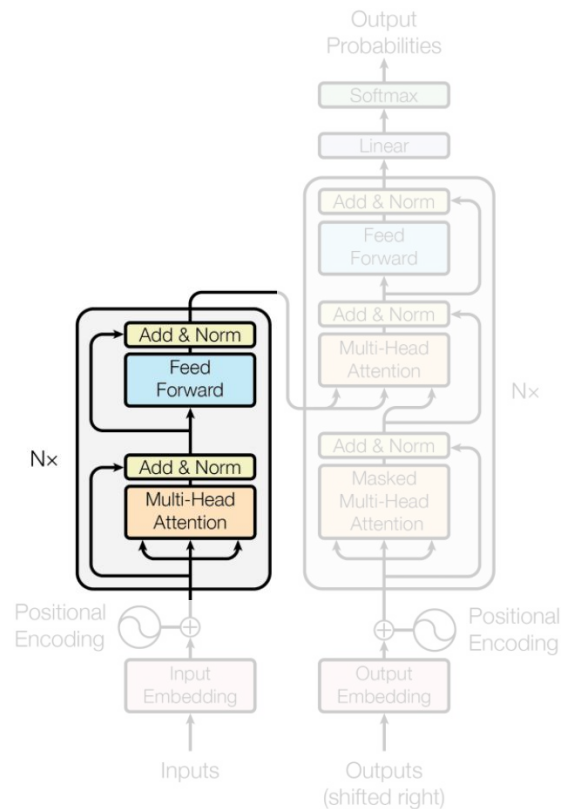
# Building blocks of Transformer

$N$  blocks, each has

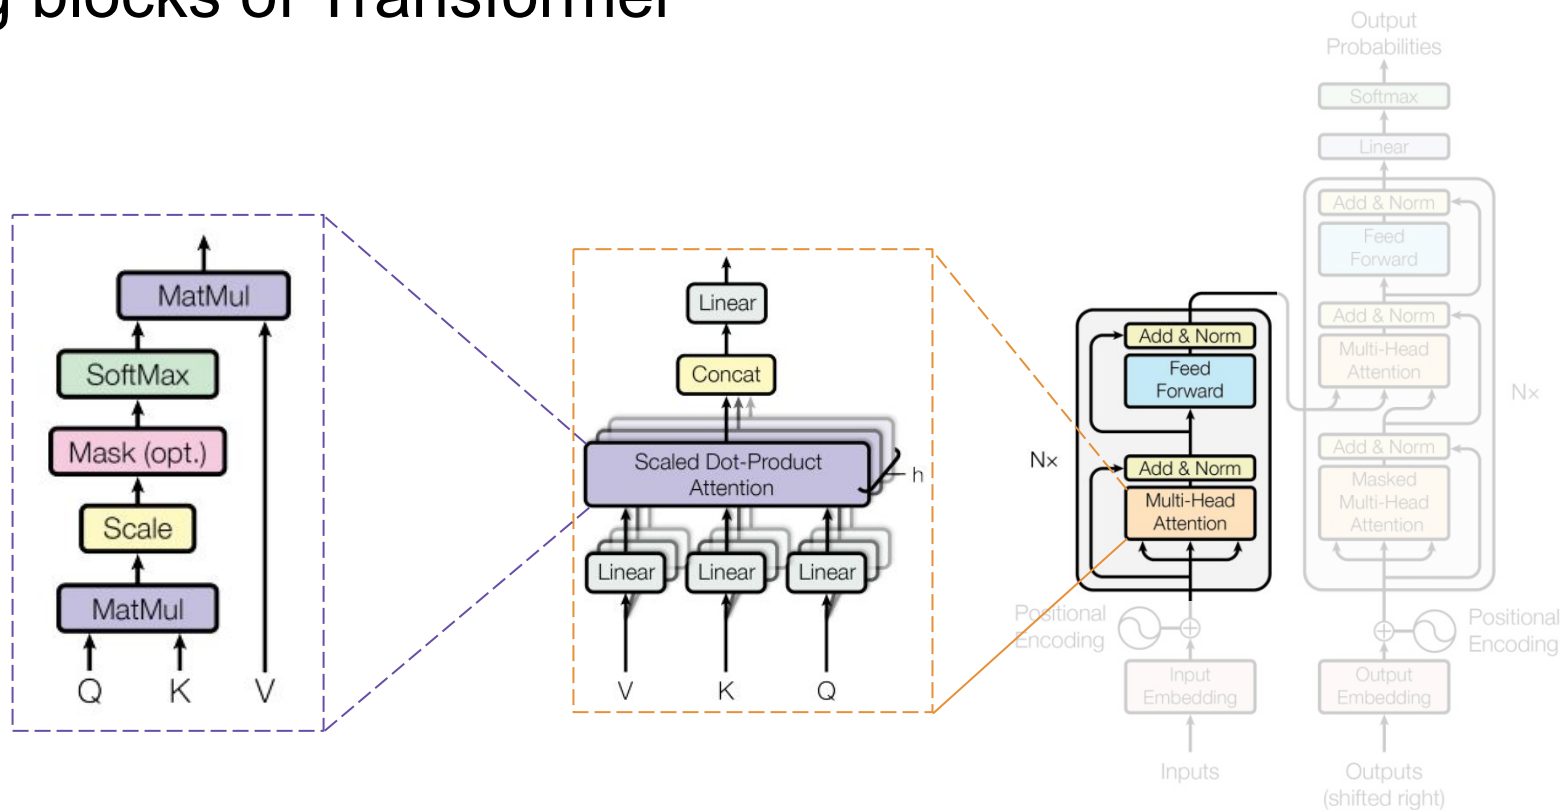
- Multi-head self-attention layer
- Two-layer feed-forward neural nets

Residual connection and layer normalization are used

- Reading: LayerNorm (<https://arxiv.org/pdf/1607.06450.pdf>)




# Building blocks of Transformer



# Key ideas: self-attention layer

- **Attention layer:** a layer to learn the dependency between words in the input. The dependency is quantified using “attention weights”
- For each word, a new representation is computed by weight-averaging the old representations of all words, where the weight is the learned attention weight




# Key ideas: self-attention layer

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

**Q**: “query” matrix, a vector representation for each word

**K**: “key” matrix, a vector representation for each word

**V**: “value” matrix, a vector representation for each word






# Key ideas: positional encoding


- Self-attention does not know the order of input words
- Positional encodings are added to the word representations, so same words at different locations have different overall representations

$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$
$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$



# Summary of Transformer

- Learning temporal relationships without unrolling and without RNNs
- Encoder/Decoder framework, multi-head self-attention modules
- Widely used in state-of-the-art NLP models
- Readings:
  - “Attention is all you need” (<https://arxiv.org/abs/1706.03762>)
  - PyTorch [implementation](#) and [tutorial](#) of Transformer



# Demo: DL for sequence data

[Google Colab](#)

# Exercise

- The model in the Colab Notebook was implemented in Tensorflow Keras. As an exercise, re-implement the model in PyTorch
- The model in the Colab Notebook was a CNN. Implement a different neural network (RNN, Transformer), then train and test it.

# Overview of Large Language Models (LLMs)


# Outline

- **What is LLM?**
- Model architectures & (Pre-)Training
- From language modeling to ChatGPT (next lecture)
- LLM for biological sequence (next lecture)

# Language modeling in natural language

- **Language modeling** is the task of predicting what word comes next

*“The students opened their \_\_\_\_\_”*



- More formally: given a sequence of words  $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(t)}$ , compute the probability distribution of the next word  $\mathbf{x}^{(t+1)}$ :

$$P(\mathbf{x}^{(t+1)} \mid \mathbf{x}^{(t)}, \dots, \mathbf{x}^{(1)})$$

where  $\mathbf{x}^{(t+1)}$  can be any word in the vocabulary  $V = \{\mathbf{w}_1, \dots, \mathbf{w}_{|V|}\}$ .

- A system that does this is called a **language model (LM)**

# Language Modeling

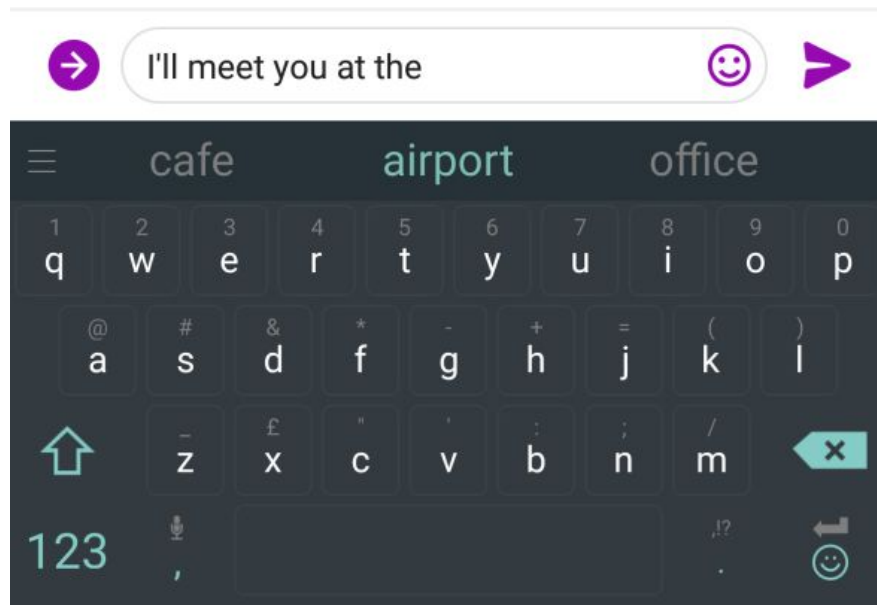
- You can also think of a Language Model as a system that **assigns a probability to a piece of text**
- For example, if we have some text  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(T)}$ , then the probability of this text (according to the Language Model) is:

$$\begin{aligned} P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(T)}) &= P(\mathbf{x}^{(1)}) \times P(\mathbf{x}^{(2)} | \mathbf{x}^{(1)}) \times \dots \times P(\mathbf{x}^{(T)} | \mathbf{x}^{(T-1)}, \dots, \mathbf{x}^{(1)}) \\ &= \prod_{t=1}^T P(\mathbf{x}^{(t)} | \mathbf{x}^{(t-1)}, \dots, \mathbf{x}^{(1)}) \end{aligned}$$

**This is what our LM provides**




# You use LM every day!




# You use LM every day!



what is the | 

what is the **weather**  
what is the **meaning of life**  
what is the **dark web**  
what is the **xfl**  
what is the **doomsday clock**  
what is the **weather today**  
what is the **keto diet**  
what is the **american dream**  
what is the **speed of light**  
what is the **bill of rights**

# You use LM every day!



GitHub Copilot

<https://copilot.github.com/>

# n-gram Language Models

*the students opened their \_\_\_\_\_*

- **Question:** How to learn a Language Model?
- **Answer** (pre- Deep Learning): learn an *n*-gram Language Model!
- **Definition:** An *n*-gram is a chunk of *n* consecutive words.
  - *unigrams*: “the”, “students”, “opened”, “their”
  - *bigrams*: “the students”, “students opened”, “opened their”
  - *trigrams*: “the students opened”, “students opened their”
  - *four*-grams: “the students opened their”
- **Idea:** Collect statistics about how frequent different n-grams are and use these to predict next word.

# n-gram Language Models

- First we make a **Markov assumption**:  $x^{(t+1)}$  depends only on the preceding  $n-1$  words

$$P(x^{(t+1)} | x^{(t)}, \dots, x^{(1)}) = P(x^{(t+1)} | \overbrace{x^{(t)}, \dots, x^{(t-n+2)}}^{n-1 \text{ words}})$$

(assumption)

prob of a n-gram

prob of a (n-1)-gram

$$= \frac{P(x^{(t+1)}, x^{(t)}, \dots, x^{(t-n+2)})}{P(x^{(t)}, \dots, x^{(t-n+2)})}$$

(definition of conditional prob)

- Question:** How do we get these  $n$ -gram and  $(n-1)$ -gram probabilities?
- Answer:** By **counting** them in some large corpus of text!

$$\approx \frac{\text{count}(x^{(t+1)}, x^{(t)}, \dots, x^{(t-n+2)})}{\text{count}(x^{(t)}, \dots, x^{(t-n+2)})}$$

(statistical approximation)

# n-gram Language Models: Example

Suppose we are learning a 4-gram Language Model.

~~as the proctor started the clock, the~~ students opened their \_\_\_\_\_  
discard condition on this

$$P(w | \text{students opened their}) = \frac{\text{count}(\text{students opened their } w)}{\text{count}(\text{students opened their})}$$

For example, suppose that in the corpus:


- “students opened their” occurred 1000 times
- “students opened their books” occurred 400 times
  - $\rightarrow P(\text{books} | \text{students opened their}) = 0.4$
- “students opened their exams” occurred 100 times
  - $\rightarrow P(\text{exams} | \text{students opened their}) = 0.1$

# How to build a *neural* language model?

- Recall the Language Modeling task:
  - Input: sequence of words  $x^{(1)}, x^{(2)}, \dots, x^{(t)}$
  - Output: prob. dist. of the next word  $P(x^{(t+1)} | x^{(t)}, \dots, x^{(1)})$

~~us~~   ~~the~~   ~~proctor~~   ~~started~~   ~~the~~   ~~clock~~     
                      discard

the   students   opened   their   \_\_\_\_\_  
                      fixed window



# We can mask everywhere in a sentence


- Autoregressive language models (predict the next word given preceding words)

*“The students opened their \_\_\_\_\_”*

- Masked language models (predict the masked word given surrounding words)

*“The students \_\_\_\_\_ their book”*

- Training process of language models (self-supervised training)
  - Randomly mask one or more words in a given sentence (e.g., from Wikipedia)
  - Train the LM (a neural network) to predict the correct word for the masked positions





# What can we learn from reconstructing the input?

Georgia Institute of Technology is located in \_\_\_\_\_, Georgia.

# What can we learn from reconstructing the input?

I put \_\_\_\_ fork down on the table.

# What can we learn from reconstructing the input?

The woman walked across the street,  
checking for traffic over \_\_\_\_ shoulder.

# What can we learn from reconstructing the input?

I went to the ocean to see the fish, turtles, seals, and \_\_\_\_\_.

# What can we learn from reconstructing the input?

Overall, the value I got from the two hours watching it was the  
sum total of the popcorn  
and the drink. The movie was \_\_\_\_.

# What can we learn from reconstructing the input?

I was thinking about the sequence that goes  
1, 1, 2, 3, 5, 8, 13, 21, \_\_\_\_\_

# Outline

- What is LLM?
- **Model architectures & (Pre-)Training**
- From language modeling to ChatGPT (next lecture)
- LLM for biological sequence (next lecture)


# Pretraining through language modeling

Recall the **language modeling** task:

- Model  $p_{\theta}(w_t|w_{1:t-1})$ , the probability distribution over words given their past contexts.
- There's lots of data for this! (In English.)

**Pretraining through language modeling:**

- Train a neural network to perform language modeling on a large amount of text.
- Save the network parameters.






# The Pretraining / Finetuning Paradigm

Pretraining can improve NLP applications by serving as parameter initialization.


## Step 1: Pretrain (on language modeling)

Lots of text; learn general things!



## Step 2: Finetune (on your task)


Not many labels; adapt to the task!



# Pre-training on related unlabeled data helps!

**Task:** Train an LM to generate product review

“Conventional” approach



Small-scale data (e.g., Amazon product reviews)

Likely to generate low-quality texts, with grammar or semantic errors

Pre-training & fine-tuning paradigm



Large-scale related data (e.g., Wikipedia articles)

First learn how to write in English, without grammar or semantic errors

Small-scale data (e.g., Amazon product reviews)

Then adapt the “writing skills” to write special-purpose texts


# Stochastic gradient descent and pretrain/finetune

Why should pretraining and finetuning help, from a “training neural nets” perspective?

- Consider, provides parameters  $\hat{\theta}$  by approximating  $\min_{\theta} \mathcal{L}_{\text{pretrain}}(\theta)$ .
  - (The pretraining loss.)
- Then, finetuning approximates  $\min_{\theta} \mathcal{L}_{\text{finetune}}(\theta)$ , starting at  $\hat{\theta}$ .
  - (The finetuning loss)
- The pretraining may matter because stochastic gradient descent sticks (relatively) close to  $\hat{\theta}$  during finetuning.
  - So, maybe the finetuning local minima near  $\hat{\theta}$  tend to generalize well!
  - And/or, maybe the gradients of finetuning loss near  $\hat{\theta}$  propagate nicely!


# Three types of LM architectures

The neural architecture influences the type of pretraining, and natural use cases




**Decoders**

- Language models! What we've seen so far.
- Nice to generate from; can't condition on future words



**Encoders**

- Gets bidirectional context – can condition on future!
- How do we train them to build strong representations?




**Encoder-  
Decoders**

- Good parts of decoders and encoders?
- What's the best way to pretrain them?


# Three types of LM architectures

The neural architecture influences the type of pretraining, and natural use cases




**Decoders**

- Language models! What we've seen so far.
- Nice to generate from; can't condition on future words



**Encoders**

- Gets bidirectional context – can condition on future!
- How do we train them to build strong representations?



**Encoder-  
Decoders**

- Good parts of decoders and encoders?
- What's the best way to pretrain them?

# Pretraining decoders


It's natural to pretrain decoders as language models and then use them as generators, finetuning their  $p_{\theta}(w_t|w_{1:t-1})$ !

This is helpful in tasks **where the output is a sequence** with a vocabulary like that at pretraining time!

- Dialogue (context=dialogue history)
- Summarization (context=document)

$$h_1, \dots, h_T = \text{Decoder}(w_1, \dots, w_T)$$
$$w_t \sim Ah_{t-1} + b$$

Where  $A, b$  were pretrained in the language model!



[Note how the linear layer has been pretrained.]

# Pretraining decoders


When using language model pretrained decoders, we can ignore that they were trained to model  $p(w_t|w_{1:t-1})$ .

We can finetune them by training a classifier on the last word's hidden state.

$$h_1, \dots, h_T = \text{Decoder}(w_1, \dots, w_T)$$
$$y \sim Ah_T + b$$

Where  $A$  and  $b$  are randomly initialized and specified by the downstream task.


Gradients backpropagate through the whole network.



[Note how the linear layer hasn't been pretrained and must be learned from scratch.]


# Three types of LM architectures

The neural architecture influences the type of pretraining, and natural use cases




**Decoders**

- Language models! What we've seen so far.
- Nice to generate from; can't condition on future words



**Encoders**

- Gets bidirectional context – can condition on future!
- How do we train them to build strong representations?



**Encoder-  
Decoders**

- Good parts of decoders and encoders?
- What's the best way to pretrain them?




# Pretraining encoders: what pretraining objective to use?

So far, we've looked at language model pretraining. But **encoders get bidirectional context**, so we can't do language modeling!

Idea: replace some fraction of words in the input with a special [MASK] token; predict these words.

$$h_1, \dots, h_T = \text{Encoder}(w_1, \dots, w_T)$$
$$y_i \sim Aw_i + b$$

Only add loss terms from words that are “masked out.” If  $\tilde{x}$  is the masked version of  $x$ , we're learning  $p_\theta(x|\tilde{x})$ . Called **Masked LM**.




# BERT: Bidirectional Encoder Representations from Transformers

Devlin et al., 2018 proposed the “Masked LM” objective and **released the weights of a pretrained Transformer**, a model they labeled BERT.


Some more details about Masked LM for BERT:

- Predict a random 15% of (sub)word tokens.
  - Replace input word with [MASK] 80% of the time
  - Replace input word with a random token 10% of the time
  - Leave input word unchanged 10% of the time (but still predict it!)
- Why? Doesn't let the model get complacent and not build strong representations of non-masked words. (No masks are seen at fine-tuning time!)




# Three types of LM architectures

The neural architecture influences the type of pretraining, and natural use cases




**Decoders**

- Language models! What we've seen so far.
- Nice to generate from; can't condition on future words



**Encoders**

- Gets bidirectional context – can condition on future!
- How do we train them to build strong representations?



**Encoder-  
Decoders**


- Good parts of decoders and encoders?
- What's the best way to pretrain them?

# Pretraining encoder-decoders: what pretraining objective to use?

For **encoder-decoders**, we could do something like **language modeling**, but where a prefix of every input is provided to the encoder and is not predicted.

$$\begin{aligned}h_1, \dots, h_T &= \text{Encoder}(w_1, \dots, w_T) \\h_{T+1}, \dots, h_{2T} &= \text{Decoder}(w_1, \dots, w_T, h_1, \dots, h_T) \\y_i &\sim Ah_i + b, i > T\end{aligned}$$

The **encoder** portion benefits from bidirectional context; the **decoder** portion is used to train the whole model through language modeling.




## T5: Denoising the span corruption

Replace different-length spans from the input with unique placeholders; decode out the spans that were removed!

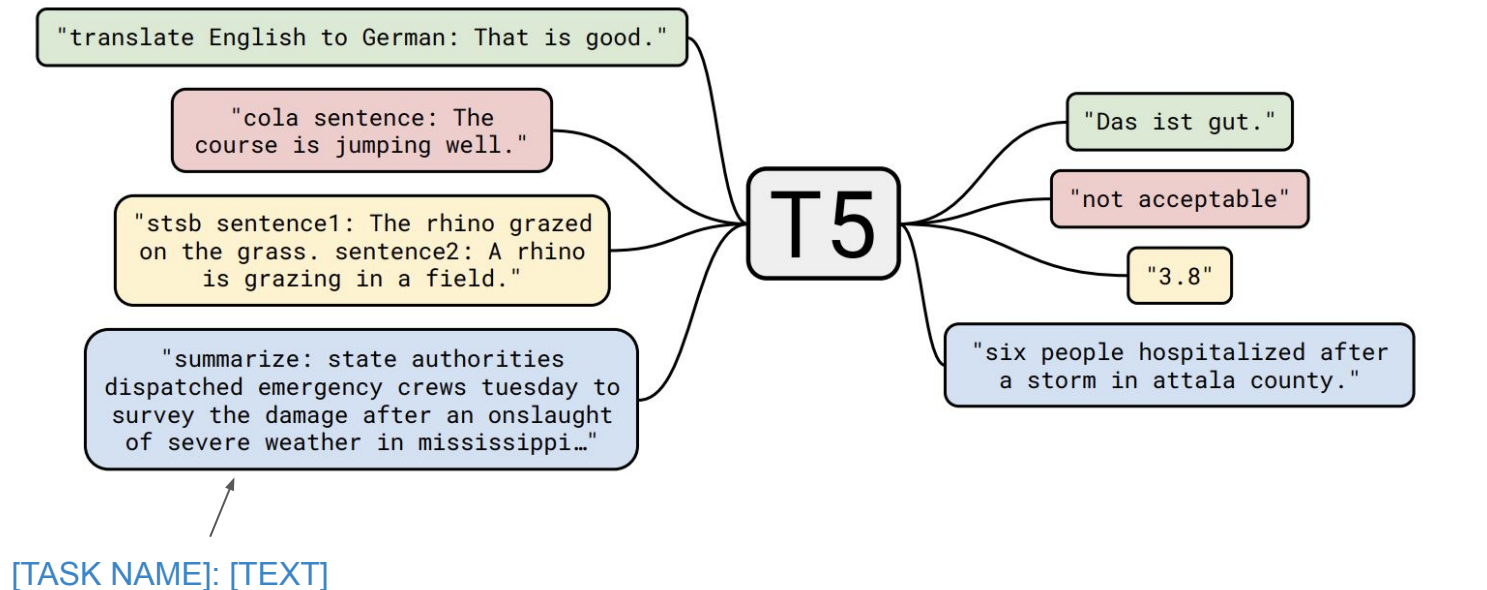
Original text

Thank you for inviting me to your party last week.

This is implemented in text preprocessing: it's still an objective that looks like **language modeling** at the decoder side.



# T5: a unified sequence-to-sequence model



# Summary of today

- Biological sequences
  - DNA, RNA, protein
- CNNs
  - Extract spatial features using convolution filters
- RNNs
  - Capture temporal dependency
  - LSTM, GRU
- Transformers
  - Self-attention layer
  - Widely used in recent state-of-the-art NLP models
- Overview of language models (LMs)
  - LM: predicting a word given context
- Next:
  - Delve into LLMs
  - How to build ChatGPT?
  - LLMs for biological sequences