

CSE7850/CX4803 Machine Learning in Computational Biology




Lecture 21: Programmable Protein Design

Yunan Luo

Controllable protein Design

- **Inverse folding**
 - Structure -> Sequence

Protein structure



Condition

Algorithm / Model



Amino acid sequence

MEKVNFLKNGVLRPPGFRFRPTDEELVVQYLKRKVFSPASPASIPEVEVYKSDPVDLPDMEQEKYFFSTK
EVKYPNGNRSNRATNSGYWKATGIDKQIILRGQQQQLIGLKKTLVIFYRGKSPHGCRTNWIMHEYRLAN
LESNYHPIQGNWVICRIFLKKRGNTKNKEENMTTHDEVNRNREIDKNPVSVMSSRDSEALASANSELKK

Design

Sometimes, our goal could not be easily specified by a 3D structure input (because we don't know such structures)


Examples:

- Generate a protein whose 99% residues are in alpha-helix
 - Or 80% residues, 50% residues, ..
- Generate a protein that can exhibit green fluorescence when exposed to light in the blue to ultraviolet range
- Generate a antibody protein that can bind to SARS-CoV-2 virus

Property-guided protein design


Gold standard is wet-lab experiment, but often use computational model for design efficiency

Generative AI model




Designed protein sequence

MEKVNFLKNGVLRPPGFRFRPTDEELVVQYLKRKVSFPLPASIIPEVEVYKSDPWDLPGDMEQEYFFSTK
EVKYPNGNRSNRATNSGYWKATGIDKQIILRGRQQQLIGLKKTLVFYRGKSPHGCRTNWIMHEYRLAN
LESNYHPIQGNVVICRIFLKKRGNTKNKEENMTTHDEVNRNREIDKNSPVSVKMSRDSEALASANSELKK



Property evaluator



- Generate a protein whose 99% residues are in alpha-helix
 - Count the residues
- Generate a protein that can exhibit green fluorescence when exposed to light in the blue to ultraviolet range
 - An ML model to predict the fluorescence brightness
- Generate a antibody protein that can bind to SARS-CoV-2 virus
 - An ML model to predict the binding affinity

Today's papers

A high-level programming language for generative protein design

Brian Hie^{1,2*} Salvatore Candido^{1*} Zeming Lin^{1,3} Ori Kabeli¹
Roshan Rao¹ Nikita Smetanin¹ Tom Sercu¹ Alexander Rives^{1,4†}

Paper #1:

Use **gradient-free** scoring function to guide the design

Article

Illuminating protein space with a programmable generative model

<https://doi.org/10.1038/s41586-023-06728-8>

Received: 20 December 2022

Accepted: 6 October 2023


Published online: 15 November 2023

John B. Ingraham¹, Max Baranov¹, Zak Costello¹, Karl W. Barber¹, Wujie Wang¹, Ahmed Ismail¹, Vincent Frappier¹, Dana M. Lord¹, Christopher Ng-Thow-Hing¹, Erik R. Van Vlack¹, Shan Tie¹, Vincent Xue¹, Sarah C. Cowles¹, Alan Leung¹, João V. Rodrigues¹, Claudio L. Morales-Perez¹, Alex M. Ayoub¹, Robin Green¹, Katherine Puentes¹, Frank Oplinger¹, Nishant V. Panwar¹, Fritz Obermeyer¹, Adam R. Root¹, Andrew L. Beam¹, Frank J. Poelwijk¹ & Gevorg Grigoryan^{1,2‡}


Paper #2:

Use property evaluator (ML models) for which we can compute **gradients** to guide the design


Paper #1




B 3-fold
(AAA)



4-fold
(AAAA)




5-fold
(AAAAA)




6-fold
(AAAAAA)




7-fold
(AAAAAAA)



8-fold
(AAAAA AAA)



A Hierarchical constraints
E.g., symmetric dimer of 2-fold symmetry




Tetramer of 2-fold
((AA) (AA) (AA) (AA))



Tetramer of 3-fold
((AAA) (AAA) (AAA) (AAA))




Tetramer of 4-fold
((AAAA) (AAAA) (AAAA) (AAAA))




Paper #1

Goal: generate a alpha-helix protein



Mixed alpha/beta



Full alpha




%(alpha)

$$p_{\theta}(\mathbf{x}) = \frac{\exp(-E_{\theta}(\mathbf{x}))}{Z(\theta)}$$

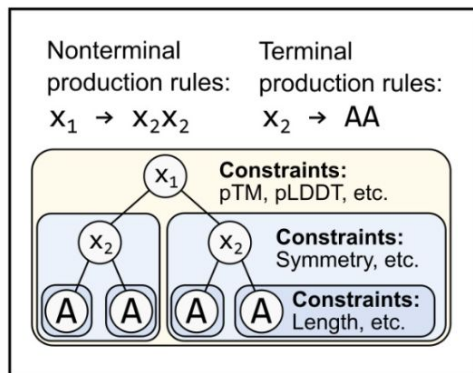
Sequence with high contents of alpha-helix will be sampled high probability

Question: how to define $E(x)$ if we want to design a symmetric protein?



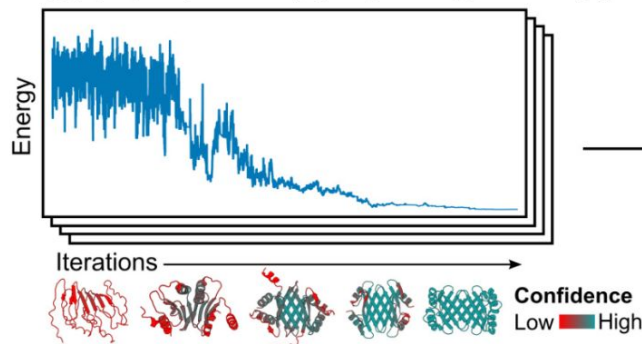
Paper #1

A Specify the design in the high-level programming language



B Define an energy function and optimize different designs


$$\text{Energy}(x) = \text{pTM}(\text{ESMFold}(x)) + \text{symmetry}(\text{ESMFold}(x)) + \dots$$



C Obtain and evaluate designed proteins

A → MSTVQ... A → RVDTN...

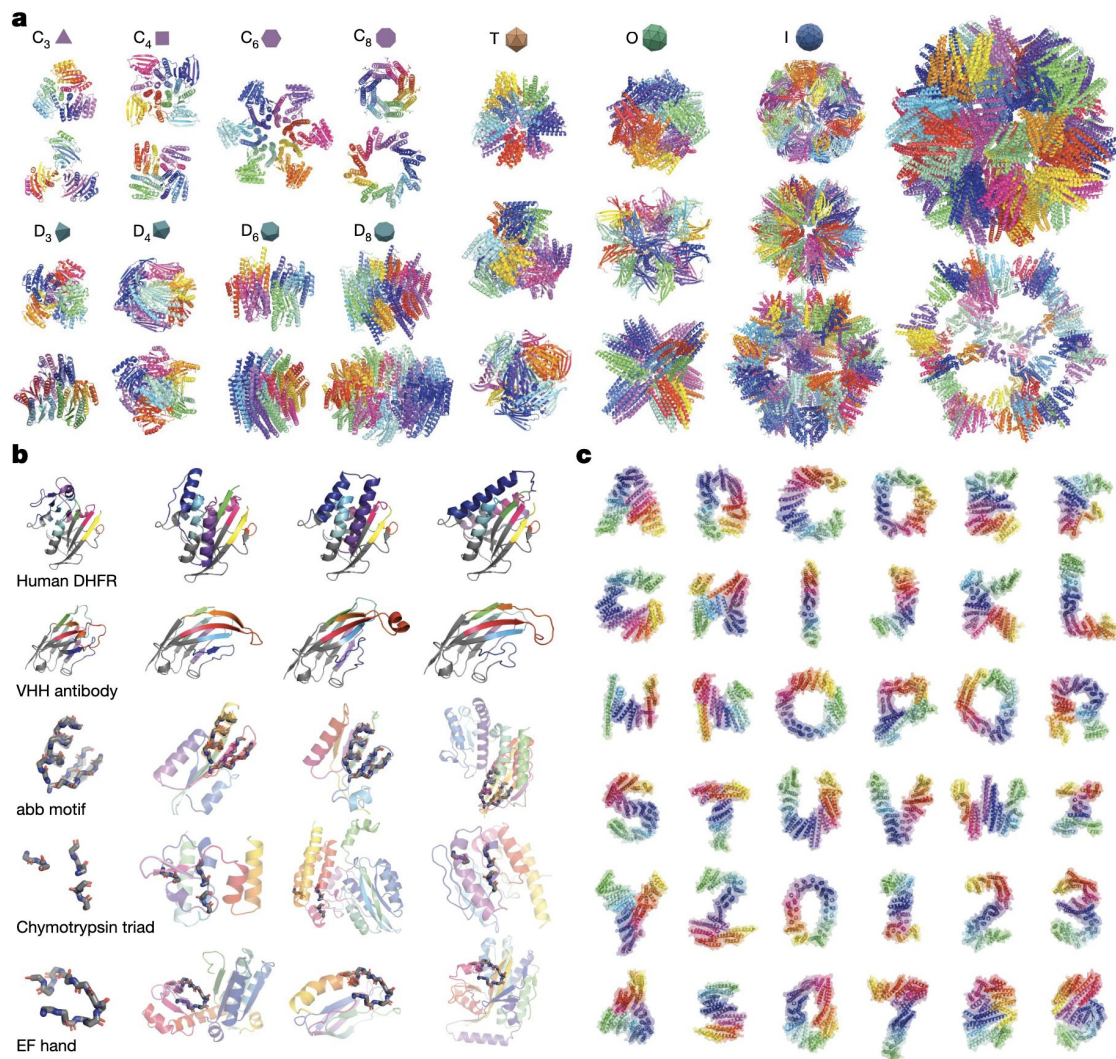

A → SAVTL... A → NLQTS...




Program — Compiles to —> Energy function — Generates —> Amino-acid sequence

Paper #2

Condition on secondary structure




Paper #2




$p(\text{protein})$

Paper #2



$$p(\text{protein}|\text{function}) \quad p(\text{protein})$$

Paper #2



$$p(\text{protein}|\text{function}) \propto p(\text{protein}) \times p(\text{function}|\text{protein})$$

Conditional
generation

Unconditional
generation

An ML
predictor!

$$\begin{aligned} \nabla_{\mathbf{x}} \log p_t(\mathbf{x}|\mathbf{y}) &= \nabla_{\mathbf{x}} \log \frac{p_t(\mathbf{x})p_t(\mathbf{y}|\mathbf{x})}{p_t(\mathbf{y})} \\ &= \nabla_{\mathbf{x}} \log p_t(\mathbf{x}) + \nabla_{\mathbf{x}} \log p_t(\mathbf{y}|\mathbf{x}) - \cancel{\nabla_{\mathbf{x}} \log p_t(\mathbf{y})} \\ &= \nabla_{\mathbf{x}} \log p_t(\mathbf{x}) + \nabla_{\mathbf{x}} \log p_t(\mathbf{y}|\mathbf{x}) \end{aligned}$$

Unconditional
generative model

Property
predictor

Diffusion Models

Ho et al. Denoising diffusion probabilistic models (DDPM), Neurips 2020.

Song et al. Score-based generative modeling through stochastic differential equations, ICLR 2021.

Bao et al. Analytic-DPM: an Analytic Estimate of the Optimal Reverse Variance in Diffusion Probabilistic Models, ICLR 2022.

Bao et al. Estimating the Optimal Covariance with Imperfect Mean in Diffusion Probabilistic Models, ICML 2022.

Rombach et al. High-resolution image synthesis with latent diffusion models. CVPR, 2022.

Text-to-image generation

Input

An astronaut riding a horse in photorealistic style.


Output



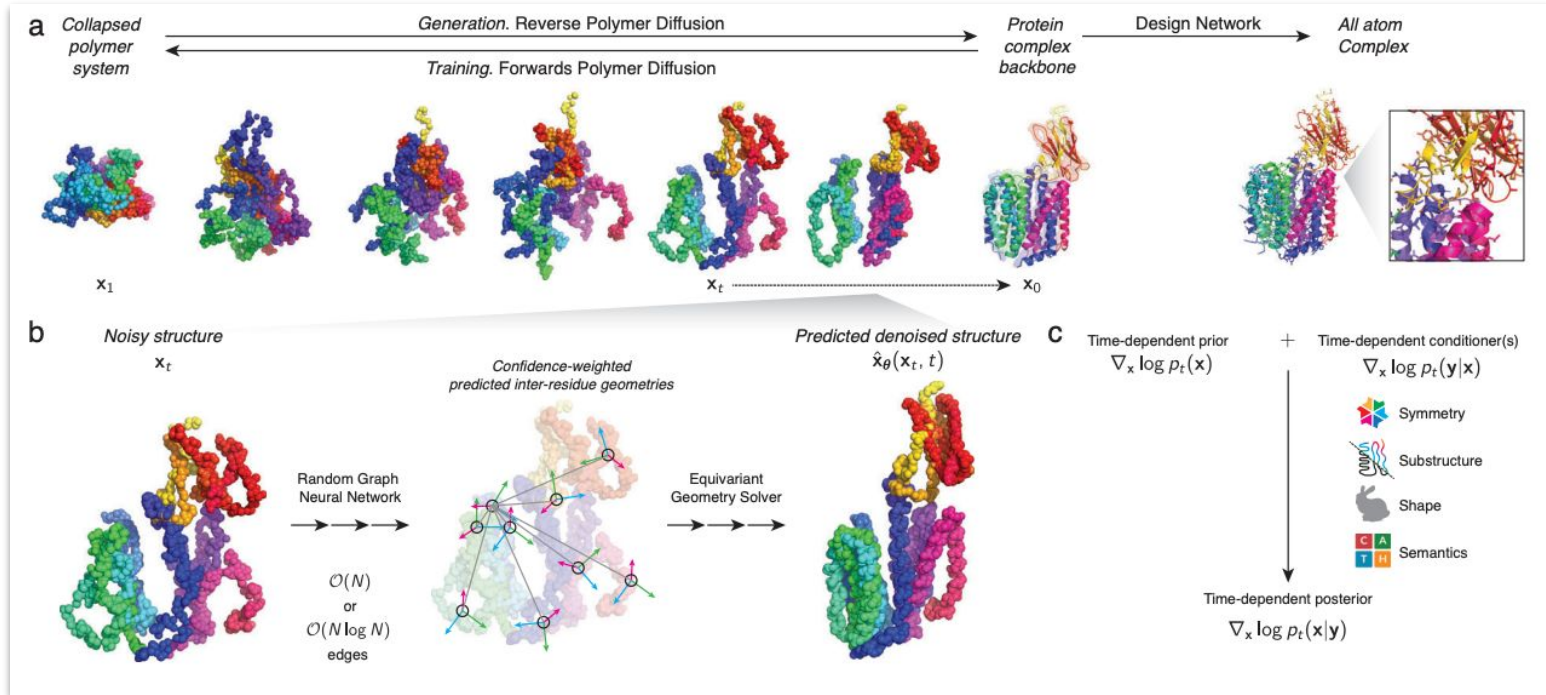
Diffusion models

Denoising diffusion models consist of two processes:


- Forward diffusion process that gradually adds noise to input
- Reverse denoising process that learns to generate data by denoising



Application: Protein Design



Application: Drug Design

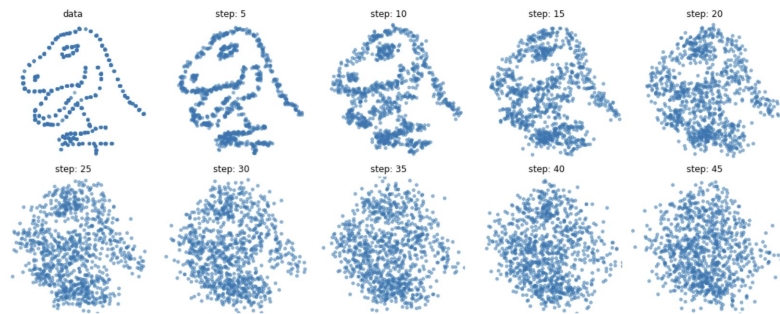


Simple code demo

<https://github.com/tanelp/tiny-diffusion>

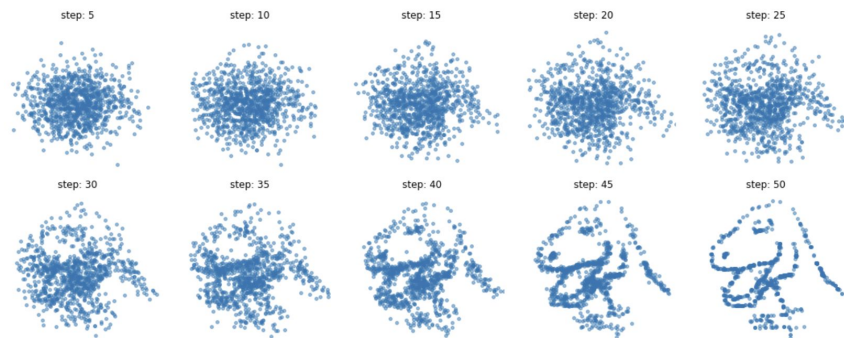
Forward process

A visualization of the forward diffusion process being applied to a dataset of one thousand 2D points. Note that the dinosaur is not a single training example, it represents each 2D point in the dataset.



Reverse process

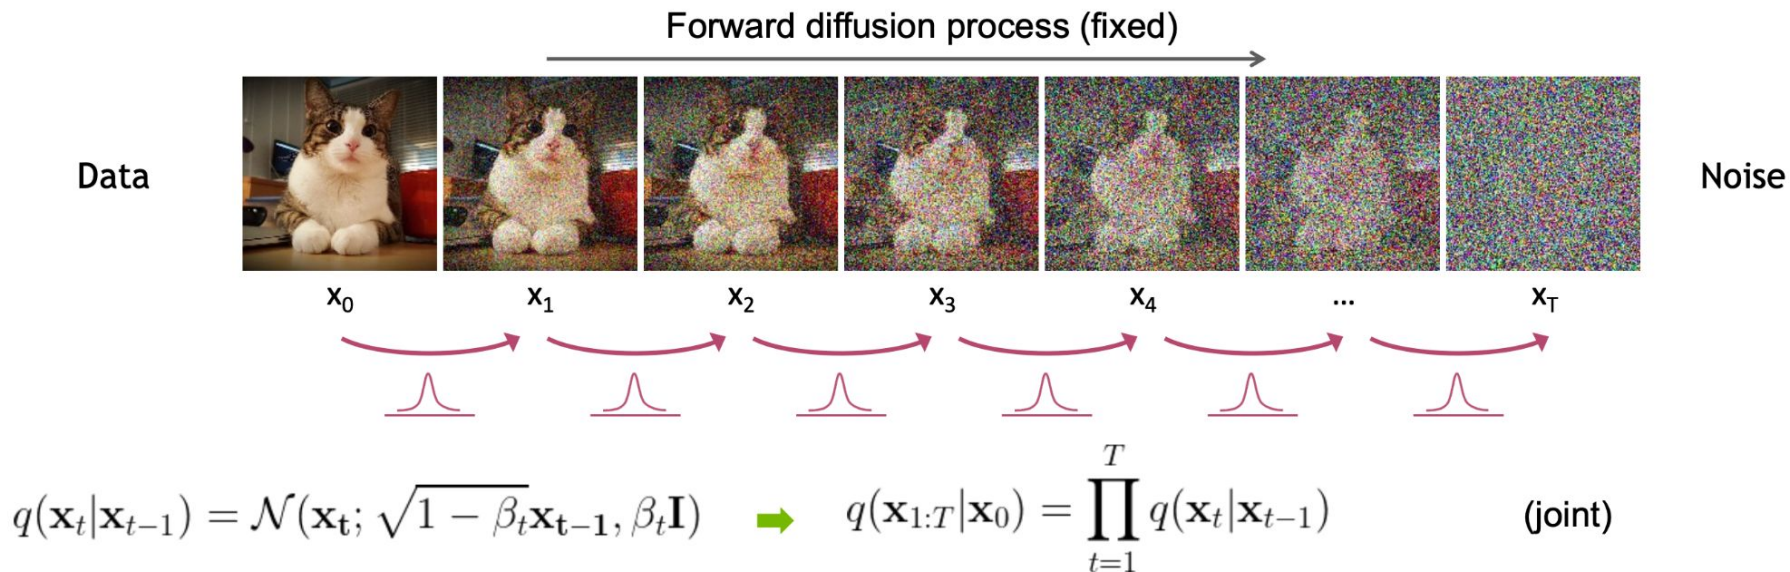
This illustration shows how the reverse process recovers the distribution of the training data.



Forward Diffusion Process

```
def add_noise(self, x_start, x_noise, timesteps):  
    s1 = self.sqrt_alphas_cumprod[timesteps]  
    s2 = self.sqrt_one_minus_alphas_cumprod[timesteps]  
  
    s1 = s1.reshape(-1, 1)  
    s2 = s2.reshape(-1, 1)  
  
    return s1 * x_start + s2 * x_noise
```

The formal definition of the forward process in T steps:



Reverse Denoising Process

Formal definition of forward and reverse processes in T steps:

```
def step(self, model_output, timestep, sample):  
    t = timestep  
    pred_original_sample = self.reconstruct_x0(sample, t, model_output)  
    pred_prev_sample = self.q_posterior(pred_original_sample, sample, t)  
  
    variance = 0  
    if t > 0:  
        noise = torch.randn_like(model_output)  
        variance = (self.get_variance(t) ** 0.5) * noise  
  
    pred_prev_sample = pred_prev_sample + variance  
  
    return pred_prev_sample
```

